

Ethernet Packet Sniffer

Introduction

An Ethernet Packet Sniffer is a program used to capture packets being sent over a network. It is a very useful networking tool that can be used in debugging networks, testing application communication and used for such questionable activities like monitoring or spying on user traffic and reverse engineering proprietary protocols.

This report documents the procedures in Lab 1 in which an Ethernet packet sniffer program was developed using the C programming language. The requirements for the program were given as follows:

- The program should print out packets in a suitable format.
- The program must capture and identify at least the following
 - IPv4 traffic
 - TCP
 - UDP
 - ICMP
 - unknown
 - IPv6 traffic
 - IPv6 Extension headers
 - TCP
 - UDP
 - ICMPv6
 - unknown
 - other types of Ethernet traffic

As promiscuous mode was unavailable on the ECS Arch Linux machines, tcpdump was used to capture the packet data, capturing data was therefore excluded from the requirements listed.

The scope of this document is to provide in depth details on how the requirements were achieved and show understanding of internet packets and protocols.

Design

The Ethernet packet sniffer made use of the libpcap framework which is a system independent interface to capture and identify packets. This includes all the C header files which contained the packet structures, header type values and the offline packet processing loop method. The packet processing method pulls out one Ethernet packet for each loop iteration so that it is possible to process what each packet contains.

The requirements stated that the internet protocol/network layer packet had to be determined and then if possible the transport layer protocol.

The IEEE 802.3 standard for Ethernet states that an Ethernet header should conform to the design in figure 1.

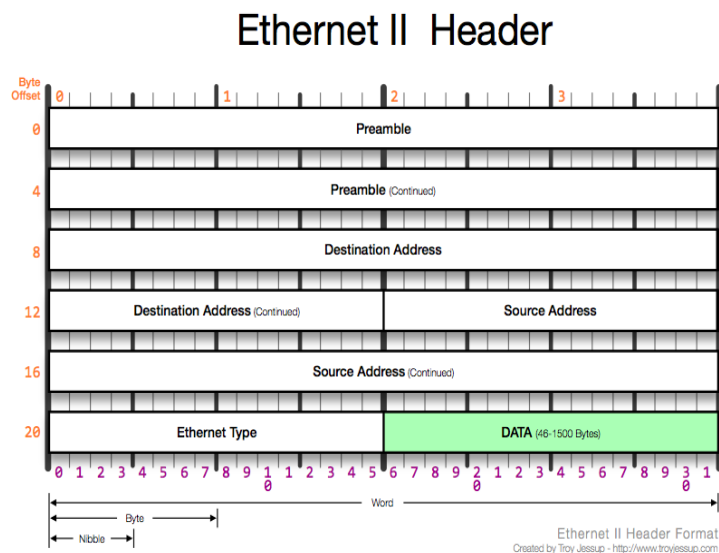


Figure 1: Ethernet II Header

As shown in this diagram it is possible to see that an Ethernet header contains an Ethernet type which is a 16 bit/2 byte value. To get this value we just have to reference it from the libpcap Ethernet packet structure. Once the Ethernet type value is retrieved, it is converted from network byte order to host byte order and the type is compared to the values for IPv4, IPv6 and the Address Resolution Protocol (ARP). All other values are being treated as unknown for this implementation.

IPv4

If it has been determined that the header after the Ethernet header is an IPv4 header, the C structure must be retrieved. It is done by typecasting the location of the packet location plus the size of the Ethernet header to get the correct place in memory where the IP header is stored.

The IPv4 header shown in figure 2 as described in the Internet Engineering Task Force (IETF) publication RFC 791.

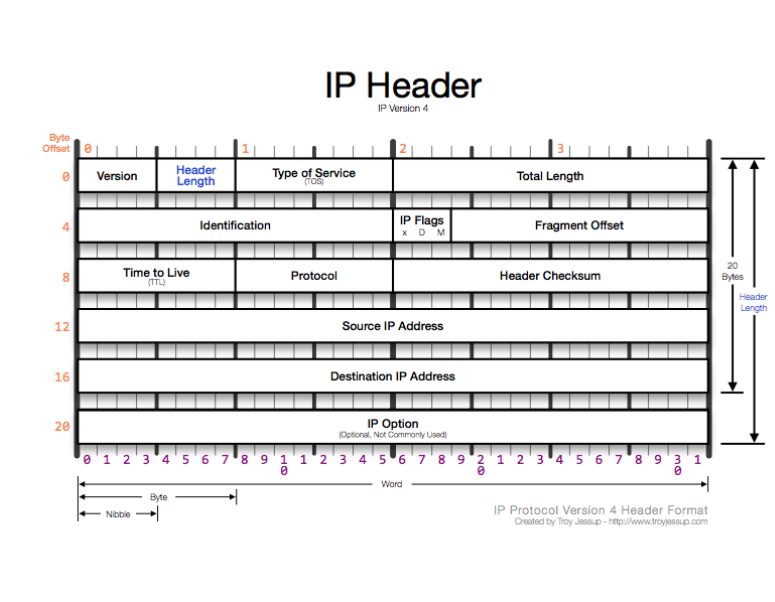


Figure 2: IPv4 Header

As seen in figure 2, the IPv4 header contains some useful information that the program could potentially print out. The packet sniffer implemented however only printed the source and destination address. When getting the source and destination address from the IPv4 header the C function `inet_ntop` had to be used to convert the 32 bit addresses into dotted decimal human readable format.

The next layer we have to find is the transport layer protocol. In IPv4 the header contains a value that defines what the transport layer protocol is. In the C code, the protocol can just be referenced from the `libpcap` IPv4 header structure. The protocol value is then compared to the values for TCP, UDP and ICMP which will determine what the next header will be.

TCP & UDP

If the next value found to be TCP then the TCP C structure must be retrieved. If it is UDP then the UDP structure is retrieved. This process is similar to the getting the IP layer protocol except we also add the size of the last header.

The TCP header shown in figure 3 as described in the IETF publication RFC 793 and the UDP header in figure 4 as described in the IETF publication RFC 786.

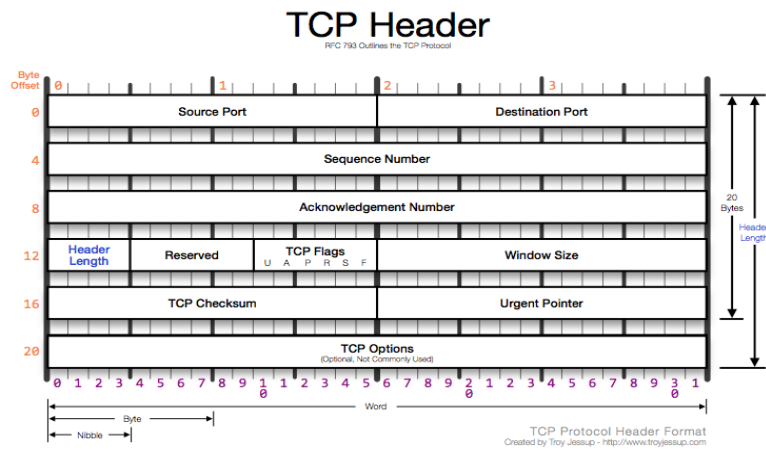


Figure 3: TCP Header

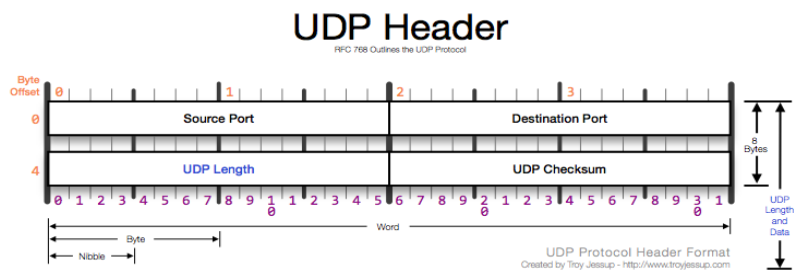


Figure 4: UDP Header

The source and destination ports, checksum, UDP and TCP payload were printed out. These values apart from the payload had to be converted to host byte order as explained in IPv4. To print the payload, the TCP and UDP header size had to be calculated so all data contained in the rest of the packet was considered to be payload. The data was casted to an unsigned char and the pointer passed to the PrintData function referenced from binarytides.com. This function converts any unreadable characters to human readable format and prints the data.

ICMP

The ICMP header shown in figure 5 as described in the IETF publication RFC 792.

The ICMP process is very similar to the TCP and UDP process except that it doesn't contain as much information. As we can see from the ICMP header in figure 5 it we can get the type, code and checksum from the packet. We just get the values from the C libpcap structure and convert the values to host byte order and get the payload as everything after the header.




Figure 5: ICMP Header

IPv6

Once we have determined that the header after the Ethernet header is an IPv6 header we can extract the information we need and determine the transport layer protocol contained inside.

The IPv6 header shown in figure 6 as described in the IETF publication RFC 2460.




Figure 6: IPv6 Header

From looking at this header we can determine a lot of information. In the application, the source and destination addresses were printed but the C function `inet_ntop` had to be used to convert the 128 bit addresses to a human readable hexadecimal format. IPv6 differs from IPv4 as it doesn't have to have a transport protocol straight after it. IPv6 can have any amount of extension headers after it, where the first one's type is given in the Next Header field shown in figure 6. Every extension header after this can have a next header field until the transport layer protocol is found.

A recursive function was used to determine each next header as the extension headers were recursed through. The function called `IPV6_HANDLER` would take the packet pointer, the point where we are up to in the packet (size of all the headers

so far), and the next header as arguments along with a string array to store values for printing tidily later on in the program. On the first function call the next header value from the IPv6 header would be passed as an argument so the function would determine what the next header is by matching the next headers value with the value for the next extension header or transport layer protocol.

Once determined, if the value is an extension header, the C structure of the header would be retrieved like previous headers and the next header value, size of all previous headers including the size of this extension header, packet and string will be passed to the same function for a recursive call.

If the next header value is a transport layer protocol then the same situation applies as the IPv4 transport layer protocols.

ARP

If the header after the Ethernet header is found to be an ARP header then we retrieve the ARP header structure in C. From the header in figure 7 which is defined by RFC 826 we can determine the opcode that tells us the operation of the ARP packet. Next the hardware type is used to then determine the Internet protocol type. In the program Ethernet 10/100Mbps was the hardware type implemented and then IPv4 was the Internet protocol implemented under it. Since we know the internet protocol is IPv4 the source and destination can now be printed as an IPv4 address.

Hardware type		Protocol type
HW addr lth	P addr lth	Opcode
Source hardware address		
Source protocol address		
Destination hardware address		
Destination protocol address		

ARP message

Figure 7: ARP Header

Apparatus and Procedures

The equipment and tools need to replicate this experiment is as follows:

- A computer running an operating system that supports:
 - o libpcap or windows version called winpcap.
 - o C compiler like GCC or MinGW for windows
 - o Wireshark or Equivalent.
 - o TCPDUMP or Equivalent.
- Experience in C programming and an understanding of Internet layers.

Results and Discussion

The results were verified by checking packets from the programs output with Wireshark a pre-existing packet sniffer application. All information printed matched the Wireshark output. This means that the Ethernet packet sniffer was functioning correctly. The results of the testing are displayed in Appendix 2.

Conclusions and recommendations.

The Ethernet packet sniffer was a success as it passed all the tests and met all the requirements of the lab. It shows understanding of how the libpcap library and Internet protocols work. Future work could contain more work on implementing other protocols and building on the ARP implementation as currently it only detects Ethernet IPv4 packets.

References

TCPDUMP & LIBPCAP (<http://www.tcpdump.org>)

TCPDUMP and LIBPCAP are used under a 3-clause BSD license.

ARP IETF RFC - <http://tools.ietf.org/html/rfc826>

TCP IETF RFC - <http://tools.ietf.org/html/rfc793>

UDP IETF RFC - <http://tools.ietf.org/html/rfc786>

ETHERNET IEEE 802.3 - <http://standards.ieee.org/about/get/802/802.3.html>

IPV6 IETF RFC - <http://tools.ietf.org/html/rfc2460>

ICMPv6 IETF RFC - <http://tools.ietf.org/html/rfc4443>

ICMP IETF RFC - <http://tools.ietf.org/html/rfc792>

Figures from - <http://www.troyjessup.com/headers/> provided as a public resource.

Appendices

Appendix A – Source Code

```
/* NWEN302 LAB 1
 * Name: Adam Bates
 * Usercode: batesadam
 * Student ID: 300223031
 */

// Libraries/Header Files to include
#include </usr/include/netinet/ip.h>
#include </usr/include/netinet/ip6.h>
#include </usr/include/pcap/pcap.h>
#include <stdio.h>
#include <stdlib.h>
#include <errno.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <netinet/if_ether.h>
#include <net/ethernet.h>
#include <netinet/ether.h>
#include <netinet/tcp.h>
#include <netinet/udp.h>
#include <netinet/ip_icmp.h>
#include <netinet/icmp6.h>
#include <stdbool.h>
#include <string.h>

/* Function Prototypes */
void packetHandler(u_char *, const struct pcap_pkthdr*, const u_char*);
void IPV6_HANDLER(int, int, const u_char*, char*);
void Handle_TCP (const u_char*, int*);
void Handle_UDP (const u_char*, int*);
void PrintData (const u_char *, int Size);
void Handle_ARP(const u_char*, int*);
void printIPv4header(char*, char*);
void Handle_ICMPV6(const u_char*, int*);
void printIPv6Header();

/* Global Variables */
int packet_counter = 0, p = 0;
bool ARP_bool = false, IPV4_bool = false, IPV6_bool = false, ICMP_bool =
false;
bool TCP_bool = false; UDP_bool = false; UNKNOWN_bool = false;
UNKNWONPROTO_bool = false;

/* Declare IPv6 Source and Destination Address Variables */
char sourceIp6[INET6_ADDRSTRLEN];
char destIp6[INET6_ADDRSTRLEN];

int main(int argc, char *argv[]) {
```

```

/* Filename of PCAP file as the first argument of the program
   if it is not included then inform the user and terminate program */
if(argc == 1){
    printf("Please include pcap file as first argument\n");
    return;
}

const char* filename = argv[1];

int i;
for(i=2; i < argc; i++){
    if(strcasecmp("ARP", argv[i]) == 0)    // strcmp ignores case when
    comparing strings
        ARP_bool = true;
    else if(strcasecmp("TCP", argv[i]) == 0)
        TCP_bool = true;
    else if(strcasecmp("UDP", argv[i]) == 0)
        UDP_bool = true;
    else if(strcasecmp("IPV4", argv[i]) == 0)
        IPV4_bool = true;
    else if(strcasecmp("IPV6", argv[i]) == 0)
        IPV6_bool = true;
    else if(strcasecmp("ICMP", argv[i]) == 0)
        ICMP_bool = true;
    else if(strcasecmp("UNKNOWN", argv[i]) == 0)
        UNKNOWN_bool = true;
}

if(argc == 2){ /* If there are no filters then enable all */
    ARP_bool = true; IPV4_bool = true; IPV6_bool = true; UNKNOWN_bool =
true;
}

    if((IPV4_bool == true || IPV6_bool == true) && TCP_bool == false &&
UDP_bool == false && ICMP_bool == false && UNKNWONPROTO_bool == false){
        TCP_bool = true; UDP_bool = true; ICMP_bool = true; UNKNWONPROTO_bool
= true;
    }

    // Select a protocol filter
    char errbuf[PCAP_ERRBUF_SIZE];

    // Open capture file
    pcap_t *descr;
    descr = pcap_open_offline(filename, errbuf);
    if(descr == NULL){
        printf("Name of file was: %s\n", filename);
        printf("Error, : %s\n", errbuf);
    }

    // start packet processing loop, just like live capture
    if (pcap_loop(descr, 0, packetHandler, NULL) < 0) {
        //cout << "pcap_loop() failed: " << pcap_geterr(descr);
        printf("pcap_loop() failed\n");
        return 1;
    }
}

```

```

else{
    printf("Loop not failed\n");
}

return 1;
}

void packetHandler(u_char *userData, const struct pcap_pkthdr* pkthdr, const
u_char* packet) {

    /* Link Layer - Declear Ethernet Header */
    const struct ether_header* ethernet_header;

    /* Declear IPv4 Headers */
    const struct ip* ip_header;
    const struct udphdr* udp_header;
    const struct icmphdr* icmp_header;

    /* Declear Source and Destination IPv4 Address Variables */
    char sourceIp[INET_ADDRSTRLEN];
    char destIp[INET_ADDRSTRLEN];

    /* Declear IPv6 Headers */
    const struct ip6_hdr* ip6_header;

    p = pkthdr->len;

    packet_counter++; // Increment the packet counter

    /* Initialise Ethernet Structure */
    ethernet_header = (struct ether_header*)packet;

    int size = 0;
    size+=sizeof(struct ether_header);

    switch(ntohs(ethernet_header->ether_type)){
    case ETHERTYPE_IP: // IPV4 Header

        if(IPV4_bool == false) return;

        ip_header = (struct ip*)(packet + size);
        inet_ntop(AF_INET, &(amp;ip_header->ip_src), sourceIp, INET_ADDRSTRLEN);
        inet_ntop(AF_INET, &(amp;ip_header->ip_dst), destIp, INET_ADDRSTRLEN);

        size+=sizeof(struct ip);

        u_char *data;
        int dataLength = 0;

        switch(ip_header->ip_p){
        case IPPROTO_TCP: // Transmission Control Protocol (TCP)
            if(TCP_bool == false) return;
            printIPV4header(sourceIp, destIp);
            Handle_TCP(packet, &size);
            break;
        case IPPROTO_UDP:
            if(UDP_bool == false) return;

```

```

        printIPV4header(sourceIp, destIp);
        Handle_UDP(packet, &size);
        break;
    case IPPROTO_ICMP: // Internet Control essgae Protocol (ICMP)

        if(ICMP_bool == false) return;
        printIPV4header(sourceIp, destIp);
        printf(" Protocol: ICMP\n");

        icmp_header = (struct icmphdr*)(packet + sizeof(struct
ether_header) + sizeof(struct ip));

        u_int type = icmp_header->type;

        if(type == 11){
            printf(" TTL Expired\n");
        }
        else if(type == ICMP_ECHOREPLY){
            printf(" ICMP Echo Reply\n");
        }

        data = (u_char*)(packet + sizeof(struct ether_header) +
sizeof(struct ip) + sizeof(struct icmphdr));
        dataLength = pkthdr->len - (sizeof(struct ether_header) +
sizeof(struct ip) + sizeof(struct icmphdr));

        printf(" Code: %d\n", (unsigned int)(icmp_header->code));
        printf(" Checksum: %d\n", ntohs(icmp_header->checksum));
        printf(" Payload(%d bytes):\n", dataLength);

        PrintData(data, dataLength);

        break;
    default: // Unknown IPV4 Protocol
        if(UNKNWONPROTO_bool == false) return;
        printf(" Protocol: Unknown\n");
        break;
    }
    break;
    case ETHERTYPE_IPV6: // IPV6

        if(IPV6_bool == false) return;

        ip6_header = (struct ip6_hdr*)(packet + size);

        inet_ntop(AF_INET6, &(ip6_header->ip6_src), sourceIp6,
INET6_ADDRSTRLEN);
        inet_ntop(AF_INET6, &(ip6_header->ip6_dst), destIp6,
INET6_ADDRSTRLEN);

        int nexthdr = ip6_header->ip6_nxt;

        size+=sizeof(struct ip6_hdr);

        char string[100] = " ";

        IPV6_HANDLER(nexthdr, size, packet, string);

```

```

        break;
    case ETHERTYPE_ARP:    // ARP
        if(ARP_bool == false) return;
        Handle_ARP(packet, &size);
        break;
    default:
        if(UNKNOWN_bool == false) return;
        printf(" ETHER_TYPE: Unknown\n");
        break;
}

}

/* Handle IPV6 Headers */
void IPV6_HANDLER(int hrd, int size, const u_char* packet, char* string){

    switch(hrd){
    case IPPROTO_ROUTING: /* Routing Header */
        strcat(string, "ROUTING, ");
        struct ip6_rthdr* header = (struct ip6_rthdr*)(packet + size);
        size+=sizeof(struct ip6_rthdr);
        IPV6_HANDLER(header->ip6r_nxt, size, packet, string);
        break;
    case IPPROTO_HOPOPTS: /* Hop-by-Hop options */
        strcat(string, "HOP-BY_HOP, ");
        struct ip6_hbh* header_hop = (struct ip6_hbh*)(packet + size);
        size+=sizeof(struct ip6_hbh);
        IPV6_HANDLER(header_hop->ip6h_nxt, size, packet, string);
        break;
    case IPPROTO_FRAGMENT: /* Fragmentation header(FRAGMENT) */
        strcat(string, "FRAGMENTATION, ");
        struct ip6_frag* header_frag = (struct ip6_frag*)(packet + size);
        size+=sizeof(struct ip6_frag);
        IPV6_HANDLER(header_frag->ip6f_nxt, size, packet, string);
        break;
    case IPPROTO_DSTOPTS: /* Destination options(DSTOPTS) */
        strcat(string, "Destination options, ");
        struct ip6_dest* header_dest = (struct ip6_dest*)(packet + size);
        size+=sizeof(struct ip6_dest);
        IPV6_HANDLER(header_dest->ip6d_nxt, size, packet, string);
        break;
    case IPPROTO_TCP:      /* TCP PROTOCOL */
        if(TCP_bool == false) return;
        printIPV6Header();
        printf("%s\n", string);
        Handle_TCP (packet, &size);
        break;
    case IPPROTO_UDP:      /* UDP PROTOCOL */
        if(UDP_bool == false) return;
        printIPV6Header();
        printf("%s\n", string);
        Handle_UDP (packet, &size);
        break;
    case IPPROTO_ICMPV6:   /* ICMP6*/
        if(ICMP_bool == false) return;

```

```

        printIPV6Header();
        printf("%s\n", string);
        Handle_ICMPV6(packet, &size);
        break;
    default:
        if(UNKNWONPROTO_bool == false) return;
        printIPV6Header();
        printf("Unknown header(%d)", hrd); /* Unknown Header */
        break;
    }
}

void printIPV6Header(){
    printf("\n*****\n");
    printf("Packet Number: %d\n IP_Version: IPV6\n Source IP: %s\n Destination IP: %s\n Extension Headers:", packet_counter, sourceIp6, destIp6);
}

void Handle_ICMPV6(const u_char* packet, int* size){
    printf("\n");
    printf(" Protocol: ICMP\n");
    u_char *data;
    int dataLength = 0;

    struct icmp6_hdr* header_icmp6 = (struct icmp6_hdr*)(packet+*size);

    data = (u_char*)(packet + *size + sizeof(struct icmp6_hdr));
    dataLength = p - *size + sizeof(struct icmp6_hdr);

    printf(" Payload(%d bytes):\n", dataLength);

    PrintData(data, dataLength);
}

/* Handle ARP Headers */
void Handle_ARP(const u_char* packet, int* size){

    const struct ether_arp* arp_header;
    arp_header = (struct ether_arp*)(packet+*size);

    printf("\n*****\n");
    printf("Packet Number: %d\n", packet_counter);

    /* Determine the ARP Operation Type */
    printf("ARP Operation: ");
    switch(ntohs(arp_header->arp_op)){
        case ARPOP_REQUEST:
            printf("ARP Request");
            break;
        case ARPOP_REPLY:
            printf("ARP Reply");
            break;
        case ARPOP_RREQUEST:

```

```

        printf("RARP Request");
        break;
    case ARPOP_RREPLY:
        printf("RARP RARP Reply");
        break;
    case ARPOP_InREQUEST:
        printf("InARP Request");
        break;
    case ARPOP_InREPLY:
        printf("InARP Request");
        break;
    case ARPOP_NAK:
        printf("(ATM) ARP NAK");
        break;
    default:
        printf("Unknown");
        break;
}

char sourceIp[INET_ADDRSTRLEN];
char destIp[INET_ADDRSTRLEN];

/* Determine the protocol hardware identifier */
printf("\nProtocol Hardware Identifier: ");
switch(ntohs(arp_header->arp_hrd)){
    case ARPHRD_NETROM:
        printf("From KA9Q: NET/ROM pseudo");
        break;
    case ARPHRD_IEEE1394:
        printf("IEEE 1394 IPv4 - RFC 2734");
        break;
    case ARPHRD_SLIP:
        printf("Serial Line Internet Protocol (SLIP)");
        break;
    case ARPHRD_ETHER:
        printf("Ethernet 10/100Mbps.");

        /* Determine the Protocol Type */
        printf("\nProtocol: ");
        int i;
        switch(ntohs(arp_header->arp_pro)){
            case ETHERTYPE_IP:
                printf("IPv4\n");

                printf("Sender MAC: ");

                for(i=0; i<6;i++)
                    printf("%02X:", arp_header->arp_sha[i]);

                printf("\nSender IP: ");

                inet_ntop(AF_INET, &(arp_header->arp_spa), sourceIp,
INET_ADDRSTRLEN);
                printf("%s", sourceIp);

                printf("\nDestination MAC: ");

```

```

        for(i=0; i<6;i++)
            printf("%02X:", arp_header->arp_tha[i]);

        printf("\nDestination IP: ");

        inet_ntop(AF_INET, &(arp_header->arp_tpa), destIp,
INET_ADDRSTRLEN);

        printf("%s", destIp);

        printf("\n");

        break;
    default:
        printf("Unknown");
        break;
    }
    break;
case ARPHRD_APPLETALK:
    printf("APPLEtalk");
    break;
case ARPHRD_IEEE802:
    printf("IEEE 802.2 Ethernet/TR/TB");
    break;
default:
    printf("Unknown");
    break;
}
}
/* Function to handle TCP Headers */
void Handle_TCP (const u_char* packet, int* size){

    /* Initialise TCP header structure */
    const struct tcphdr* tcp_header;
    u_int sourcePort, destPort;
    u_char *data;

    tcp_header = (struct tcphdr*)(packet + *size);
    int dataLength = 0;

    /* Get the source and destination ports from the TCP header */
    sourcePort = ntohs(tcp_header->source);
    destPort = ntohs(tcp_header->dest);

    /* Initialise the data pointer to point to the data carried by
the TCP and Initialise dataLength to the length of the data */
    *size+=tcp_header->doff*4;
    data = (u_char*)(packet + *size);
    dataLength = p - *size;

    /* Print the TCP header information and payload */
    printf(" Protocol: TCP\n      Source Port: %d\n      Destination Port:
%d\n Checksum: %d\n Payload(%d bytes):\n",
        sourcePort, destPort, ntohs(tcp_header->check),
dataLength);

    /* Print the packet contents */
    PrintData (data , dataLength);

```



```

}

/* Function to handle UDP Headers */
void Handle_UDP (const u_char* packet, int* size){

    /* Initialise UDP header structure */
    const struct udphdr* udp_header;
    u_int sourcePort, destPort;
    u_char *data;

    udp_header = (struct udphdr*)(packet + *size);
    int dataLength = 0;

    /* Get the source and destination ports from the UDP header */
    sourcePort = ntohs(udp_header->source);
    destPort = ntohs(udp_header->dest);

    /* Initialise the data pointer to point to the data carried by
the UDP and Initialise dataLength to the length of the data */
    *size+=sizeof(struct udphdr);
    data = (u_char*)(packet + *size);
    dataLength = p - *size;

    /* Print the TCP header information and payload */
    printf(" Protocol: UDP\n      Source Port: %d\n      Destination Port:
%d\n Payload(%d bytes):\n",
           sourcePort, destPort, dataLength);

    /* Print the packet contents */
    PrintData (data , dataLength);
}

void printIPv4header(char* source, char* dest){

printf("*****\n");
printf("Packet number: %d\n", packet_counter);
printf(" IP version: IPv4\n");

printf("      Source IP: %s\n", source);
printf("      Destination IP: %s\n", dest);
}

/* Convert and Print Data from protocols
 * (PrintData Adapted from BinaryTides, "http://www.binarytides.com/packet-
sniffer-code-c-libpcap-linux-sockets/"
 * written by Silver Moon)
 */
void PrintData (const u_char * data , int Size)
{
    int i , j;
    for(i=0 ; i < Size ; i++)
    {
        if( i!=0 && i%16==0)    //if one line of hex printing is complete...
        {

```

```

printf("          ");
for(j=i-16 ; j<i ; j++)
{
    if(data[j]>=32 && data[j]<=128)
        printf("%c", (unsigned char)data[j]); //if its a number or
alphabet

        else printf("."); //otherwise print a dot
    }
    printf("\n");
}

if(i%16==0) printf("          ");
printf(" %02X", (unsigned int)data[i]);

if( i==Size-1) //print the last spaces
{
    for(j=0; j<15-i%16; j++)
    {
        printf("    "); //extra spaces
    }

    printf("          ");

    for(j=i-i%16 ; j<=i ; j++)
    {
        if(data[j]>=32 && data[j]<=128)
        {
            printf("%c", (unsigned char)data[j]);
        }
        else
        {
            printf(".");
        }
    }
    printf("\n\n");
}
}
}

```

0x10	00 34 00 00 00 00 3f 11 66 6f 82 c3 06 11 82 c3	.4...?. ko.....
0x20	04 b3 08 01 03 ff 00 20 90 47 00 09 53 12 00 00G.S....
0x30	00 01 00 00 00 00 00 00 00 00 00 00 00 00 00	[REDACTED]
0x40	00 00	[REDACTED]

IPv4- TCP

320 3.366812 130.195.4.179 130.195.6.8 NFS 206 V3 READ Call (Reply in 321), FH:0x855b5604 Offset:1961984 Len:4096

Frame 320: 206 bytes on wire (1648 bits), 206 bytes captured (1648 bits)

Ethernet II, Src: 86:2b:2b:b2:be:7d (86:2b:2b:b2:be:7d), Dst: Cisco_5c:16:7f (60:73:5c:5c:16:7f)

Internet Protocol Version 4, Src: 130.195.4.179 (130.195.4.179), Dst: 130.195.6.8 (130.195.6.8)

Version: 4
Header length: 20 bytes

Differentiated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))

Total Length: 192
Identification: 0x8552 (34130)

Flags: 0x02 (Don't Fragment)
Fragment offset: 0
Time to live: 64
Protocol: TCP (6)

Header checksum: 0xa4a4 [correct]
Source: 130.195.4.179 (130.195.4.179)
Destination: 130.195.6.8 (130.195.6.8)

Transmission Control Protocol, Src Port: 738 (738), Dst Port: nfs (2049), Seq: 15825, Ack: 66989, Len: 140

Source port: 738 (738)
Destination port: nfs (2049)
[Stream index: 0]
Sequence number: 15825 (relative sequence number)
[Next sequence number: 15965 (relative sequence number)]
Acknowledgment number: 66989 (relative ack number)
Header length: 32 bytes

Flags: 0x018 (PSH, ACK)
Window size value: 3833
[Calculated window size: 3833]
[Window size scaling factor: -1 (unknown)]

Checksum: 0x10f4 [validation disabled]

Payload (140 bytes):

```

0000 60 73 5c 5c 16 7f 86 2b 2b b2 be 7d 08 00 45 00 's\...+ +..E.
0010 00 c0 85 52 40 00 04 06 a4 a4 82 c3 04 b3 82 c3 ...R@.@. ....
0020 06 08 02 e2 08 01 38 8f 2d b2 1b 24 f9 0b 80 18 .....8. -.$....
0030 0e f9 10 f4 00 00 01 01 08 0a 4b ba 3f f8 00 00 .....K.?...
0040 0a 4a 80 00 00 88 63 38 20 3b 00 00 00 00 00 00 .J....c8 ;.....
0050 00 02 00 01 86 a3 00 00 00 03 00 00 00 06 00 00 .....
0060 00 01 00 00 00 34 01 1b 12 fc 00 00 00 14 67 6f .....4.. ....go
0070 74 68 61 6d 2e 65 63 73 2e 76 75 77 2e 61 63 2e tham.ecs .vuw.ac.
0080 6e 7a 00 00 9c 5f 00 00 00 19 00 00 00 03 00 00 nz.....
0090 00 19 00 00 04 41 00 00 05 c0 00 00 00 00 00 00 ....A. ....
00a0 00 00 00 00 00 1c 01 04 00 80 8b 07 00 00 0c 00 ..V...l.....
00b0 00 00 a2 af 76 01 bf 0a b3 6c 00 00 00 00 00 00 ....V...l.....
00c0 00 00 00 00 00 00 00 1d f0 00 00 00 10 00 .....

```

Packet number: 320
IP version: IPv4
Source IP: 130.195.4.179
Destination IP: 130.195.6.8
Protocol: TCP
Source Port: 738
Destination Port: 2049
Checksum: 4340
Payload(140 bytes):

```

80 00 00 00 63 38 20 3b 00 00 00 00 00 00 00 02 ?...c8 ;.....
00 01 86 a3 00 00 00 03 00 00 00 06 00 00 00 01 .....
00 00 00 34 01 1b 12 fc 00 00 00 14 67 6f 74 68 ...4.....goth
61 60 2e 65 63 73 2e 76 75 77 2e 61 63 2e 6e 7a am.ecs.vuw.ac.nz
00 00 9c 5f 00 00 00 19 00 00 00 03 00 00 00 19 .....
00 00 04 41 00 00 05 c0 00 00 00 00 00 00 00 00 ....A.....
00 00 00 1c 01 04 00 80 8b 07 00 00 0c 00 00 00 .....?.....
A2 AF 76 01 BF 0A B3 6C 00 00 00 00 00 00 00 00 ..V...l.....
00 00 00 00 00 1D F0 00 00 00 10 00 .....

```

Packet number: 321
IP version: IPv4

IPv4 – ICMP

Output:

```

*****
Packet number: 5649
IP version: IPv4
Source IP: 130.195.4.179
Destination IP: 130.195.5.12
Protocol: ICMP
Code: 3
Checksum: 3373
Payload(237 bytes):
  45 00 00 ED 30 16 00 00 3F 11 3B A5 82 C3 05 0C      E...0...?..
  82 C3 04 B3 00 35 4A 8C 00 D9 8A 1D 5D A2 81 90      .....5J.....]...
  00 01 00 01 00 04 00 05 03 67 6F 02 67 6C 00      .....goo.gl.
  00 1C 00 01 C0 0C 00 1C 00 01 00 00 01 2C 00 10      .....
  24 04 68 00 40 06 08 03 00 00 00 00 00 10 08      $.h.@.....
  C0 0C 00 02 00 01 00 00 32 59 00 10 03 6E 73 33      .....2Y...ns3
  06 67 6F 6F 67 6C 65 03 63 6F 6D 00 C0 0C 00 02      .google.com.....
  00 01 00 00 32 59 00 06 03 6E 73 31 C0 44 C0 0C      ....2Y...ns1.D..
  00 02 00 01 00 00 32 59 00 06 03 6E 73 32 C0 44      .....2Y...ns2.D
  C0 0C 00 02 00 01 00 00 32 59 00 06 03 6E 73 34      .....2Y...ns4
  C0 44 C0 5C 00 01 00 01 00 05 28 41 00 04 D8 EF      .D.\.....(A....
  20 0A C0 6E 00 01 00 01 00 05 28 41 00 04 D8 EF      ..n.....(A....
  22 0A C0 40 00 01 00 01 00 05 28 41 00 04 D8 EF      ".@.....(A....
  24 0A C0 80 00 01 00 01 00 05 28 41 00 04 D8 EF      $.?.....(A....
  26 0A 00 00 29 10 00 00 00 80 00 00 00      &...)....?...
*****

```

Wireshark:

Wireshark packet capture showing an ICMP Echo (ping) request and response between 130.195.4.179 and 130.195.5.12.

Packet 5649: 36.105513 130.195.4.179 130.195.5.12 ICMP 27

Version: 4
Header Length: 20 bytes
Differenziated Services Field: 0xc0 (DSCP 0x30: Class Selector 6; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))
1100 00.. = Differenziated Services Codepoint: Class Selector 6 (0x30)
.... ..00 = Explicit Congestion Notification: Not-ECT (Not ECN-Capable Transport) (0x00)
Total Length: 265
Identification: 0x90e2 (39138)
Flags: 0x00
Fragment offset: 0
Time to live: 64
Protocol: ICMP (1)
Header checksum: 0xd10c [correct]
Source: 130.195.4.179 (130.195.4.179)
Destination: 130.195.5.12 (130.195.5.12)

Internet Control Message Protocol
Type: 3 (Destination unreachable)
Code: 3 (Port unreachable)
Checksum: 0x0d2d [correct]

Internet Protocol Version 4, Src: 130.195.5.12 (130.195.5.12), Dst: 130.195.4.179 (130.195.4.179)
Version: 4
Header Length: 20 bytes
Differenziated Services Field: 0x00 (DSCP 0x00: Default; ECN: 0x00: Not-ECT (Not ECN-Capable Transport))
0000 00.. = Differenziated Services Codepoint: Default (0x00)
.... ..00 = Explicit Congestion Notification: Not-ECT (Not ECN-Capable Transport) (0x00)
Total Length: 237
Identification: 0x3016 (12310)
Flags: 0x00
Fragment offset: 0
Time to live: 63
Protocol: UDP (17)
Header checksum: 0x3ba5 [correct]
Source: 130.195.5.12 (130.195.5.12)
Destination: 130.195.4.179 (130.195.4.179)

User Datagram Protocol, Src Port: domain (53), Dst Port: 19084 (19084)
Domain Name System (response)
[Request In: 5597]
[Time: 0.174311000 seconds]
Transaction ID: 0x5da2
Flags: 0x8190 Standard query response, No error
Questions: 1
Answer RRs: 1

0010 01 09 98 e2 00 00 40 01 d1 0c 82 c3 04 b3 82 c3@.....
0020 05 0c 03 03 0d 2d 00 00 00 00 82 c3 04 b3 82 c30.....
0030 00 00 3f 11 3b a5 82 c3 05 0c 82 c3 04 b3 00 35?.....5
0040 4a 8c 00 d9 8a 1d 5d a2 81 90 00 01 00 01 00 04 J.....].....
0050 00 05 03 67 6f 6f 02 67 6c 00 00 1c 00 01 c0 0c ...goo.g l.....
0060 00 1c 00 01 00 00 01 2c 00 10 24 04 68 00 40 06 \$.h.@.....
0070 08 03 00 00 00 00 00 19 08 c0 0c 00 02 00 012Y...n s3.goog
0080 00 00 32 59 00 10 03 6e 73 33 06 67 6f 67 6c ..2Y...n s3.goog
0090 65 03 63 6f 6d 00 c0 0c 00 02 00 01 00 00 32 59 e.com... ..2Y
00a0 00 06 03 6e 73 31 c0 44 c0 0c 00 02 00 01 00 00 ...ns1.D.....
00b0 32 59 00 06 03 6e 73 32 c0 44 c0 0c 00 02 00 01 2Y...ns2.D.....
00c0 00 00 32 59 00 06 03 6e 73 34 c0 44 c0 5c 00 01 ..2Y...n s4.D.\..
00d0 00 01 00 05 28 41 00 04 d8 ef 20 0a c0 6e 00 01(A.....n..
00e0 00 01 00 05 28 41 00 04 d8 ef 22 0a c0 40 00 01(A.....".@..
00f0 00 01 00 05 28 41 00 04 d8 ef 24 0a c0 80 00 01(A.....\$.
0100 00 01 00 05 28 41 00 04 d8 ef 26 0a 00 00 29 10(A.....&...
0110 00 00 00 80 00 00 00</p>
</div>

IPv6 - TCP

```

9711 39.753735 2404:2000:2000:4:842b:2bff:feb2:be7d 2404:6800:4008:c00::7d TCP 86 35084 > xmpp-client [ACK] Seq=1 Ack=1 Win=258 Len=0 TSval=1270508189 TSecr=3..
> Frame 9711: 86 bytes on wire (688 bits), 86 bytes captured (688 bits)
> Ethernet II, Src: 86:2b:2b:b2:be:7d (86:2b:2b:b2:be:7d), Dst: Cisco_5c:16:7f (60:73:5c:5c:16:7f)
> Internet Protocol Version 6, Src: 2404:2000:2000:4:842b:2bff:feb2:be7d (2404:2000:2000:4:842b:2bff:feb2:be7d), Dst: 2404:6800:4008:c00::7d (2404:6800:4008:c00::7d)
  > 0110 .... = Version: 6
  > .... 0000 0000 .... = Traffic class: 0x00000000
    .... 0000 0000 0000 0000 0000 = FlowLabel: 0x00000000
    Payload length: 32
    Next header: TCP (6)
    Hop limit: 64
    Source: 2404:2000:2000:4:842b:2bff:feb2:be7d (2404:2000:2000:4:842b:2bff:feb2:be7d)
    Destination: 2404:6800:4008:c00::7d (2404:6800:4008:c00::7d)
> Transmission Control Protocol, Src Port: 35084 (35084), Dst Port: xmpp-client (5222), Seq: 1, Ack: 1, Len: 0
  Source port: 35084 (35084)
  Destination port: xmpp-client (5222)
  [Stream index: 89]
  Sequence number: 1 (relative sequence number)
  Acknowledgment number: 1 (relative ack number)
  Header length: 32 bytes
  > Flags: 0x010 (ACK)
  Window size value: 258
  [Calculated window size: 258]
  [Window size scaling factor: -1 (unknown)]
  > Checksum: 0xaa13 [validation disabled]
  > Options: (12 bytes), No-Operation (NOP), No-Operation (NOP), Timestamps

```

```

0000 60 73 5c 5c 16 7f 86 2b 2b b2 be 7d 86 dd 60 00  `s\\...+ +..}...`
0010 00 00 00 20 06 40 24 04 20 00 20 00 00 04 84 2b  ... .@$ . ....+
0020 2b ff fe b2 be 7d 24 04 68 00 40 08 0c 00 00 00  +....}$ . h.@....
0030 00 00 00 00 00 7d 89 0c 14 66 59 35 5a 89 b0 fb  ....}... .fY5Z...
0040 9d 35 80 10 01 02 aa 13 00 00 01 01 08 0a 4b ba  .5..... ..K.
0050 6a 9d e8 dd f7 d4  ].....

```

```

bates_299 — ssh — 123x46

Welcome to NetBSD!

barretts: [~] % ssh gotham
Last login: Thu Aug  8 18:52:25 2013 from barretts.ecs.vuw.ac.nz
gotham: [~] % cd NWEN302/LABS/LAB1
gotham: [LAB1] % ./sniffer output2.cap ipv6 tcp

*****
Packet Number: 9711
IP_Version: IPV6
  Source IP: 2404:2000:2000:4:842b:2bff:feb2:be7d
  Destination IP: 2404:6800:4008:c00::7d
Extension Headers:
Protocol: TCP
  Source Port: 35084
  Destination Port: 5222
Checksum: 43539
Payload(0 bytes):

```

IPv6 - ICMP

Output:

```

*****
Packet Number: 5090
IP_Version: IPv6
    Source IP: fe80::842b:2bff:feb2:d2b5
    Destination IP: ff02::16
Extension Headers: HOP-BY_HOP,

Protocol: ICMP
Payload(62 bytes):
    1A F4 00 00 00 02 02 00 00 00 FF 02 00 00 00 00 .....
    00 00 00 00 00 00 00 00 02 02 02 00 00 00 FF 02 .....
    00 00 00 00 00 00 00 00 00 00 01 FF B2 D2 B5 51 F0 .....Q.
    8B EE 38 D0 55 6E 51 F0 8B EE 38 D0 55 6E .....8.UnQ...8.Un
*****

```

Wireshark:

```

    type: IPv6 (0x8000)
  ▸ Internet Protocol Version 6, Src: fe80::842b:2bff:feb2:d2b5 (fe80::842b:2bff:feb2:d2b5), Dst: ff02::16 (ff02::16)
    ▸ 0110 .... = Version: 6
    ▸ .... 0000 0000 .... = Traffic class: 0x00000000
    .... 0000 0000 0000 0000 0000 = Flowlabel: 0x00000000
    Payload length: 56
    Next header: IPv6 hop-by-hop option (0)
    Hop limit: 1
    Source: fe80::842b:2bff:feb2:d2b5 (fe80::842b:2bff:feb2:d2b5)
    Destination: ff02::16 (ff02::16)
  ▾ Hop-by-Hop Option
    Next header: ICMPv6 (58)
    Length: 0 (8 bytes)
    ▸ IPv6 Option (Router Alert)
    ▸ IPv6 Option (PadN)
  ▸ Internet Control Message Protocol v6
    Type: Multicast Listener Report Message v2 (143)
    Code: 0
    Checksum: 0x1af4 [correct]
    Reserved: 0000
    Number of Multicast Address Records: 2
    ▸ Multicast Address Record Exclude: ff02::202
    ▸ Multicast Address Record Exclude: ff02::1:ff02::d2b5
  0000 33 33 00 00 00 16 86 2b 2b b2 d2 b5 86 dd 60 00 33.....+ +.....`
  0010 00 00 00 38 00 01 fe 80 00 00 00 00 00 00 84 2b ...8.... +.....+
  0020 2b ff fe b2 d2 b5 ff 02 00 00 00 00 00 00 00 00 +.....
  0030 00 00 00 00 00 16 3a 00 05 02 00 00 01 00 8f 00 .....:
  0040 1a f4 00 00 00 02 02 00 00 00 ff 02 00 00 00 00 .. .....
  0050 00 00 00 00 00 00 00 00 02 02 02 00 00 00 ff 02 .....
  0060 00 00 00 00 00 00 00 00 00 01 ff b2 d2 b5 .....

```

```

▼ Frame 15967: 60 bytes on wire (480 bits), 60 bytes captured (480 bits)
    WTAP_ENCAP: 1
    Arrival Time: Jul 25, 2013 14:23:18.804577000 NZST
    [Time shift for this packet: 0.000000000 seconds]
    Epoch Time: 1374718998.804577000 seconds
    [Time delta from previous captured frame: 0.000796000 seconds]
    [Time delta from previous displayed frame: 0.000796000 seconds]
    [Time since reference or first frame: 46.653890000 seconds]
    Frame Number: 15967
    Frame Length: 60 bytes (480 bits)
    Capture Length: 60 bytes (480 bits)
    [Frame is marked: False]
    [Frame is ignored: False]
    [Protocols in frame: eth:arp]
    [Coloring Rule Name: ARP]
    [Coloring Rule String: arp]
▼ Ethernet II, Src: Cisco_5c:16:7f (60:73:5c:5c:16:7f), Dst: 86:2b:2b:b2:d2:eb (86:2b:2b:b2:d2:eb)
    ▼ Destination: 86:2b:2b:b2:d2:eb (86:2b:2b:b2:d2:eb)
        Address: 86:2b:2b:b2:d2:eb (86:2b:2b:b2:d2:eb)
            .... .1. .... = LG bit: Locally administered address (this is NOT the factory default)
            .... ..0 ..... = IG bit: Individual address (unicast)
    ▼ Source: Cisco_5c:16:7f (60:73:5c:5c:16:7f)
        Address: Cisco_5c:16:7f (60:73:5c:5c:16:7f)
            .... ..0. .... = LG bit: Globally unique address (factory default)
            .... ...0 .... = IG bit: Individual address (unicast)
        Type: ARP (0x0806)
        Padding: 00000000000000000000000000000000000000000000
▼ Address Resolution Protocol (reply)
    Hardware type: Ethernet (1)
    Protocol type: IP (0x0800)
    Hardware size: 6
    Protocol size: 4
    Opcode: reply (2)
    Sender MAC address: Cisco_5c:16:7f (60:73:5c:5c:16:7f)
    Sender IP address: 130.195.4.129 (130.195.4.129)
    Target MAC address: 86:2b:2b:b2:d2:eb (86:2b:2b:b2:d2:eb)
    Target IP address: 130.195.4.171 (130.195.4.171)
```